

Meta learning symmetries by reparametrization

Authors: Allan Zhou, Tom Knowles, Chelsea Finn



Presented by: Sakshi B.

Motivation

- To build symmetry into NN, we need to know symmetry in advance.
- Example, CNNs are designed to be shift invariant.
- How to design NNs when symmetry is not as explicit?

This paper: Learn the symmetry structure from data, rather than hard-code it

Meta learning



$$\arg \max_{\phi} \log p(D|\phi) + \log p(\phi)$$

$$\arg \max_{\phi} \sum_i \log p(y_i|x_i, \phi) + \log p(\phi)$$

Optimization as a model for few-shot learning

- **Equivariance**

- A function is equivariant to a transformation if *transforming the input to the function is equivalent to transforming the output*
- Consider a network layer $\phi : \mathbb{R}^n \rightarrow \mathbb{R}^m$
- Assume we have two representations of a group G where $\pi_1(g)$ transforms the input vectors and $\pi_2(g)$ transforms the output vectors. The layer ϕ is G -equivariant w.r.t these transformations if:

$$\phi(\pi_1(g)\mathbf{v}) = \pi_2(g)\phi(\mathbf{v}) \quad \text{for all } g \in G, v \in \mathbb{R}^n$$

- *Invariance is a special case of equivariance*

- If $\pi_2 = \text{id} = I$ then we have:

$$\underbrace{\phi(\pi_1(g)\mathbf{v})}_{\text{Layer output on the transformed input}} = \pi_2(g)\phi(\mathbf{v}) = \underbrace{\phi(\mathbf{v})}_{\text{Layer output on the original input}}$$

This paper:


- Meta-learning can learn equivariant features from data
- Loses this equivariance information during gradient update
- Reparameterization trick: Factorize weight matrix W as product of symmetry matrix U and parameters vector v

$$\text{vec}(W) = Uv, \quad v \in \mathbb{R}^k, U \in \mathbb{R}^{mn \times k}$$

- Encodes the pattern by which the weights W will “share” the filter parameters v

Shared filter parameters

Reparameterization:

$$\mathbf{U} \bullet \mathbf{v} = \text{vec}(\mathbf{W}) \longrightarrow \mathbf{W}$$


group representation

symmetry matrix

parameters

reshape

layer weights

It is proved that this represents decoupled equivariant sharing patterns and filter parameters for all G-convolutions with finite group G

Example: Permutation equivariance

$$\mathbf{U} \bullet \mathbf{v} = \text{vec}(\mathbf{W}) \rightarrow \mathbf{W}$$

$\underbrace{\begin{pmatrix} 1 & 0 \\ 0 & -\frac{1}{2} \\ 0 & \frac{1}{2} \\ 1 & 0 \end{pmatrix}}_{\text{sharing matrix}} \underbrace{\begin{pmatrix} \text{orange} \\ \text{green} \\ \text{green} \\ \text{orange} \end{pmatrix}}_{\text{parameters}} = \underbrace{\begin{pmatrix} \text{orange} \\ \text{green} \\ \text{green} \\ \text{orange} \end{pmatrix}}_{\text{layer weights}} \xrightarrow{\text{reshape}} \begin{pmatrix} \text{orange} & \text{green} \\ \text{green} & \text{orange} \end{pmatrix}$

Meta-training algorithm


input: $\{\mathcal{T}_j\}_{j=1}^N \sim p(\mathcal{T})$: Meta-training tasks

input: $\{\mathbf{U}, \mathbf{v}\}$: Randomly initialized symmetry matrices and filters.

input: α, η : Inner and outer loop step sizes.

while *not done* **do**

```
    sample minibatch  $\{\mathcal{T}_i\}_{i=1}^n \sim \{\mathcal{T}_j\}_{j=1}^N$ ;  
    forall  $\mathcal{T}_i \in \{\mathcal{T}_i\}_{i=1}^n$  do  
         $\{\mathcal{D}_i^{tr}, \mathcal{D}_i^{val}\} \leftarrow \mathcal{T}_i$ ; // task data  
         $\boldsymbol{\delta}_i \leftarrow \nabla_{\mathbf{v}} \mathcal{L}(\mathbf{U}, \mathbf{v}, \mathcal{D}_i^{tr})$ ;  
         $\mathbf{v}' \leftarrow \mathbf{v} - \alpha \boldsymbol{\delta}_i$ ; // inner step  
        /* outer gradient */  
         $\mathbf{G}_i \leftarrow \frac{d}{d\mathbf{U}} \mathcal{L}(\mathbf{U}, \mathbf{v}', \mathcal{D}_i^{val})$ ;  
        /* outer step */  
         $\mathbf{U} \leftarrow \mathbf{U} - \eta \sum_i \mathbf{G}_i$ ;
```



- Since we aim to learn equivariances that are shared across tasks, the symmetry matrix doesn't update with the task.
- Only filter parameters update with task training data.
- Symmetry matrix updates using task's validation data.

(Meta-training)

Meta-train:

Inner loop
(train):

$$\mathbf{v}' \leftarrow \mathbf{v} - \alpha \nabla_{\mathbf{v}} \mathcal{L}(\mathbf{U}, \mathbf{v}, \mathcal{D}_i^{\text{train}})$$

Outer loop
(val):

$$\mathbf{U} \leftarrow \mathbf{U} - \eta \nabla_{\mathbf{U}} \mathcal{L}(\mathbf{U}, \mathbf{v}', \mathcal{D}_i^{\text{val}})$$

Meta-test: U is frozen and filter parameters v are updated

Experiments: Translation equivariance

Method	Synthetic Problems MSE (lower is better)					
	Small train dataset			Large train dataset		
	$k = 1$	$k = 2$	$k = 5$	$k = 1$	$k = 2$	$k = 5$
MAML-FC	$3.4 \pm .60$	$2.1 \pm .35$	$1.0 \pm .10$	$3.4 \pm .49$	$2.0 \pm .27$	$1.1 \pm .11$
MAML-LC	$2.9 \pm .53$	$1.8 \pm .24$	$.87 \pm .08$	$2.9 \pm .42$	$1.6 \pm .23$	$.89 \pm .08$
MAML-Conv	$.00 \pm .00$	$.43 \pm .09$	$.41 \pm .04$	$.00 \pm .00$	$.53 \pm .08$	$.49 \pm .04$
MTSR-FC (Ours)	$3.2 \pm .49$	$1.4 \pm .17$	$.86 \pm .06$	$.12 \pm .03$	$.07 \pm .02$	$.07 \pm .01$
MSR-Joint-FC (Ours)	$.25 \pm .16$	$.12 \pm .04$	$.21 \pm .03$	$.01 \pm .00$	$.08 \pm .02$	$.12 \pm .02$
MSR-FC (Ours)	$.07 \pm .02$	$.07 \pm .02$	$.16 \pm .02$	$.00 \pm .00$	$.05 \pm .01$	$.09 \pm .01$

- Baseline: Gradient based meta-learning (MAML)
- $k=1$ denotes translation equivariance, $k=2,5$ is less symmetry
- This algorithms is closest to MAML-Conv among all Non-Conv methods on $k=1$
- Outperforms all methods on $k=2,5$

Experiments: Rotation and flip equivariance

Rotation/Flip Equivariance MSE		
Method	Rot	Rot+Flip
MAML-Conv	.504	.507
MSR-Conv (Ours)	.004	.001

- Task data generated by passing randomly generated inputs through single-layer E(2) steerable CNN, equivariant to combinations of translations, discrete rotations, and reflections.
- Outperforms MAML significantly.

Contributions:

- Few-shot generalization on task distributions with shared underlying symmetries.
- Reduced number of task-specific parameters.

Drawbacks:

- All experimentation on synthetic data.
- Generalization in real-world missing.
- Works for discovering symmetries shared across a distribution of tasks. How about symmetries particular to a single task?

Thank you !